

STATS 769

Data Technologies Review

Paul Murrell

The University of Auckland

July 29, 2019

Overview

- This section of the course provides a refresh on basic data technologies.
- Our primary computing environment will be R.

Flat text files are the simplest data source.

- Data stored in rows, with multiple values per row.
- Delimited files have a delimiter between values.
- Fixed-width files use a constant number of characters per value.
- CSV (comma-separated-value) files use a comma as the delimiter.
- The first row may contain variable names.
- There may be several lines of metadata at the top of the file.
- Any software can read/write the file.

An example of a CSV file:

```
publicid,origintime,longitude,latitude,depth  
2014p172810,2014-03-05T22:33:37,175.93,-36.91,368.75  
2014p172742,2014-03-05T21:57:26,175.89,-40.60,22.46  
2014p172698,2014-03-05T21:33:31,175.86,-40.63,27.44
```

File Formats

An example of a fixed-width file with metadata and column headings:

VARIABLE : Mean TS from clear sky composite						
FILENAME : ISCCPMonthly_avg.nc						
FILEPATH : /usr/local/fer_data/data/						
SUBSET : 24 by 24 points (LONGITUDE-LATITUDE)						
TIME : 16-JAN-1995 00:00						
113.8W 111.2W 108.8W 106.2W 103.8W 101.2W						
27 28 29 30 31 32						
36.2N / 51:	272.7	270.9	270.9	269.7	273.2	275.6
33.8N / 50:	279.5	279.5	275.0	275.6	277.3	279.5
31.2N / 49:	284.7	284.7	281.6	281.6	280.5	282.2
28.8N / 48:	289.3	286.8	286.8	283.7	284.2	286.8
26.2N / 47:	292.2	293.2	287.8	287.8	285.8	288.8

R can easily read text files and create a data frame.

- `read.csv(file, header, skip)`
- `read.table(file, header, sep, skip)`
- `read.fwf(file, widths, header, skip)`

Binary file formats are more complex, but tend to be more efficient in terms of both speed and size.

- Data stored in arbitrarily complex forms.
- Require specific software to read/write.

File Formats

An example of a binary file:

```
2500 : 04 00 40 00 40 00 c1 01 08 00 c1 01 00 00 54 8d 01 00 eb 00 5a 00 0f 00 | ..@.@.....T....Z...
2524 : 00 f0 52 00 00 00 00 00 06 f0 18 00 00 00 00 04 00 00 02 00 00 00 01 00 | ..R.....
2548 : 00 00 01 00 00 00 01 00 00 00 01 00 00 00 33 00 0b f0 12 00 00 00 bf 00 | .....3.....
2572 : 08 00 08 00 81 01 09 00 00 08 c0 01 40 00 00 08 40 00 1e f1 10 00 00 00 | .....@...@.....
2596 : 0d 00 00 08 0c 00 00 08 17 00 00 08 f7 00 00 10 fc 00 2b 01 23 00 00 00 | .....+.#...
2620 : 23 00 00 00 08 00 00 6c 61 74 69 74 75 64 65 07 00 00 58 31 31 33 2e 38 | #.....latitude...X113.8
2644 : 57 07 00 00 58 31 31 31 2e 32 57 07 00 00 58 31 30 38 2e 38 57 07 00 00 | W...X111.2W...X108.8W...
2668 : 58 31 30 36 2e 32 57 07 00 00 58 31 30 33 2e 38 57 07 00 00 58 31 30 31 | X106.2W...X103.8W...X101
2692 : 2e 32 57 06 00 00 58 39 38 2e 38 57 06 00 00 58 39 36 2e 32 57 06 00 00 | .2W...X98.8W...X96.2W...
2716 : 58 39 33 2e 38 57 06 00 00 58 39 31 2e 32 57 05 00 00 33 36 2e 32 4e 05 | X93.8W...X91.2W...36.2N.
2740 : 00 00 33 33 2e 38 4e 05 00 00 33 31 2e 32 4e 05 00 00 32 38 2e 38 4e 05 | ..33.8N...31.2N...28.8N.
2764 : 00 00 32 36 2e 32 4e 05 00 00 32 33 2e 38 4e 05 00 00 32 31 2e 32 4e 05 | ..26.2N...23.8N...21.2N.
2788 : 00 00 31 38 2e 38 4e 05 00 00 31 36 2e 32 4e 05 00 00 31 33 2e 38 4e 05 | ..18.8N...16.2N...13.8N.
2812 : 00 00 31 31 2e 32 4e 04 00 00 38 2e 38 4e 04 00 00 36 2e 32 4e 04 00 00 | ..11.2N...8.8N...6.2N...
2836 : 33 2e 38 4e 04 00 00 31 2e 32 4e 04 00 00 31 2e 32 53 04 00 00 33 2e 38 | 3.8N...1.2N...1.2S...3.8
2860 : 53 04 00 00 36 2e 32 53 04 00 00 38 2e 38 53 05 00 00 31 31 2e 32 53 05 | S...6.2S...8.8S...11.2S.
2884 : 00 00 31 33 2e 38 53 05 00 00 31 36 2e 32 53 05 00 00 31 38 2e 38 53 05 | ..13.8S...16.2S...18.8S.
2908 : 00 00 32 31 2e 32 53 ff 00 0a 00 23 00 40 04 00 00 0c 00 00 00 63 08 15 | ..21.2S....#.@.....c..
2932 : 00 63 08 00 00 00 00 00 00 00 00 00 00 15 00 00 00 00 00 00 02 0a 00 | .C.....
2956 : 00 00 09 08 10 00 00 06 10 00 bb 0d cc 07 00 00 00 00 06 00 00 0c 00 | .....
2980 : 02 00 64 00 0f 00 02 00 01 00 11 00 02 00 00 00 10 00 08 00 | ..d.....
```


File Formats

An example of a binary file:

```
=====
4000 : 00 01 00 01 00 0f 00 1b aa 01 00 03 02 0e 00 01 00 02 00 0f 00 | .....
=====temperature
4021 : 66 66 66 66 66 ee 70 40 | 270.9
=====
4029 : 03 02 0e 00 01 00 03 00 0f 00 | .....
=====temperature
4039 : 66 66 66 66 66 ee 70 40 | 270.9
=====
4047 : bd 00 12 00 01 00 04 00 0f 00 6b a5 01 00 0f 00 e3 aa 01 00 | .....k.....
=====
4067 : 05 00 03 02 0e 00 01 00 06 00 0f 00 | .....
=====temperature
4079 : 9a 99 99 99 99 39 71 40 | 275.6
=====
4087 : bd 00 12 00 01 00 07 00 0f 00 4b b1 01 00 0f 00 4b b1 01 00 | .....K....K...
=====
4107 : 08 00 03 02 0e 00 01 00 09 00 0f 00 | .....
=====temperature
4119 : 66 66 66 66 66 6e 71 40 | 278.9
=====
4127 : 03 02 0e 00 01 00 0a 00 0f 00 | .....
=====temperature
4137 : 66 66 66 66 66 6e 71 40 | 278.9
```

Each binary format tends to need a specific package and the results are not necessarily as simple as a data frame:

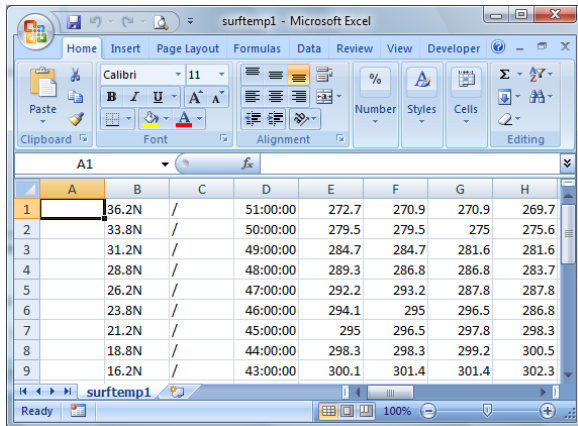
- **foreign** for a variety of statistical software binary formats (SPSS, SAS, etc)
- **RNetCDF** for netCDF files
- **hdf5** for HDF5 files

Spreadsheets are not an ideal data format, but are very common.

- Data stored in rows and columns.
- There may be multiple sheets in one workbook.
- There may be formatting for humans, which can make processing more difficult.
- Can be easily converted to flat text file (if there is only one sheet).
- Require specific software to read/write.

Spreadsheets

An example of a spreadsheet:



The screenshot shows the Microsoft Excel application window titled 'surftemp1 - Microsoft Excel'. The ribbon is set to 'Home', and the 'Clipboard' group is active. The spreadsheet contains data in columns A through H and rows 1 through 9. The data is as follows:

	A	B	C	D	E	F	G	H
1		36.2N	/	51:00:00	272.7	270.9	270.9	269.7
2		33.8N	/	50:00:00	279.5	279.5	275	275.6
3		31.2N	/	49:00:00	284.7	284.7	281.6	281.6
4		28.8N	/	48:00:00	289.3	286.8	286.8	283.7
5		26.2N	/	47:00:00	292.2	293.2	287.8	287.8
6		23.8N	/	46:00:00	294.1	295	296.5	286.8
7		21.2N	/	45:00:00	295	296.5	297.8	298.3
8		18.8N	/	44:00:00	298.3	298.3	299.2	300.5
9		16.2N	/	43:00:00	300.1	301.4	301.4	302.3

One approach with R is to convert the spreadsheet to a flat file and then use previous functions.

Otherwise, there are packages that can read spreadsheets directly, typically as data frames:

- **gdata** has `read.xls()`.
- **xlsReadWrite** (Windows only) also has `read.xls()`.
- **xlsx** has `read.xlsx()`.
- **RODBC** provides detailed access to individual spreadsheet cells.


Relational Databases

Databases provide a more sophisticated storage option, especially for large and/or complex data.

- The data are spread across multiple tables.
- Table columns are typed.
- Each table has a primary key.
- Tables are linked by foreign keys.
- Access can be limited to authorised users.
- Access can be concurrent for multiple users.
- Require very specific software to read/write.
- Standard interface via SQL.

Relational Databases

An example of a database:



Relational Databases

An example of a database:

spend_table

year	spend	spendID	ministry	minister
1972	335	1	1	2
1973	377	2	1	2
1974	440	3	1	2
1975	527	4	1	2
1976	627	5	1	1
1977	699	6	1	1
1978	808	7	1	1
1972	292	8	2	5
1973	343	9	2	3
1974	401	10	2	6
1975	492	11	2	4
1976	606	12	2	4
1977	689	13	2	4
1978	809	14	2	4

ministry_table

ministry	ministryID
Education	1
Health	2
Social	3

minister_table

minister	ministerID
Les Gandar	1
Phil Amos	2
Bob Tizard	3
Frank Gill	4
Lance Adams-Schneide	5
Tom McGuigan	6

Relational Databases

SQL can be used to access information in a database:

```
SELECT year, spend, mr.minister, my.ministry
FROM spend_table st
     INNER JOIN minister_table mr
         ON st.minister = mr.ministerID
     INNER JOIN ministry_table my
         ON st.ministry = my.ministryID
WHERE year = 1978;
```

year	spend	minister	ministry
1978	808	Les Gandar	Education
1978	809	Frank Gill	Health

There are R packages for querying databases and getting the result as a data frame:

- **ROracle**, **RMySQL**, and **RSQLite** provide `dbConnect()` and `dbGetQuery()`.

HTML and XML

HTML and XML are useful to know for several reasons: they are another potential data source format and HTML is a useful document format (for reports).

- Text files consisting of content and markup.
- Markup consists of elements and attributes.

HTML and XML

An example HTML file:

```
<html>
<head>
  <title>Electoral Donations 2014</title>
</head>
<body>
  <h1>Electoral Donations</h1>
  <p>
    Source:
    <a href="http://www.elections.org.nz/">
      The Electoral Commission of New Zealand</a>.
  </p>
</body>
</html>
```

HTML and XML

An example XML file:

```
<?xml version="1.0"?>
<ElectoralDonations>
  <party id="P2" name="National">
    <candidate id="C1" name="Amy" surname="ADAMS">
      <donation id="d1" amount="15000.00"/>
      <donation id="d2" amount="10000.00"/>
    </candidate>
  </party>
</ElectoralDonations>
```

HTML and XML

R has packages for working with HTML and XML documents. The result may be a data frame or it may be something insanely complicated:

- **XML** has `readHTMLTable()` and `xmlParse()`

We will be going into more depth with processing HTML and XML later in the course.

Data Structures

R has a small set of standard data structures:

- Vectors are 1-dimensional and have a type (character, numeric, logical).
- Matrices are 2-dimensional and have a type.
- Data frames are 2-dimensional and each column has a type.
- Lists are recursive; each component of a list can be any data structure.
- Data frame are lists where each component has to be the same length.
- Use `str()` to get a low-level view of a data structure.

Subsetting

- Single square brackets, `x[index]`, can be used to extract a range of values from `x`; the result is usually the same class as `x`.
- Double square brackets, `x[[index]]` can be used to extract a single component from a list.
- The `index` can be numeric, logical, or character.
- The syntax `x$name` is short for `x[["name"]]`.

Control Flow

R has standard conditional and loop constructs:

```
if (condition) {  
  ...  
} else {  
  ...  
}
```

```
for (i in values) {  
  ...  
}
```

```
while (condition) {  
  ...  
}
```

Writing Functions

We can define new R functions:

```
f <- function(x, y=0) {  
  ...  
}
```

- arguments can have default values.
- the last expression in the body of the function provides the return value for the function.

R has tools for manipulating data structures:

- Summary functions: `min()`, `max()`, `sum()`, `range()`, `mean()`.
- Generating sequences: `seq()`, `rep()`, `c()`.
- Tables of counts: `table()`, `ftable()`, `xtabs()`.
- Combining structures: `cbind()`, `rbind()`, `merge()`.
- Apply functions: `apply()`, `sapply()`, `lapply()`.
- Aggregation functions: `aggregate()`, `tapply()`, `by()`.
- Split-apply-recombine: `split()`, `lapply()`, `do.call()`.

Reshaping

- Long format.
- Wide format.
- The **reshape2** package.

R has tools for manipulating text:

- Search (and replace) text: `grep()`, `regexpr()`, `gsub()`.
- Break text into pieces: `strsplit()`.
- Combine text: `paste()`.
- Most functions make use of regular expressions.

Dates

- Format dates.
- Generate sequences of dates.
- Perform arithmetic on dates.

Debugging

- Use `traceback()` to show the call stack after an error.
- Use `browser()` to interrupt execution and inspect objects.
- Use `debug()` to interrupt execution when a function is called.
- Use `trace()` to interrupt execution within a function call.
- Use `recover()` to browse any currently active function calls.
- Use `options(warn)` to turn warnings into errors.
- Use `options(error)` to call `recover()` after an error.

- “Introduction to Data Technologies”
<https://www.stat.auckland.ac.nz/~paul/ItDT/>
- Ross Ihaka’s “An R Programming Quick Reference”
(on Canvas)
- Duncan Murdoch’s “Debugging in R”
<https://web.archive.org/web/20170706215053/http://www.stats.uwo.ca:80/faculty/murdoch/software/debuggingR/>